

# tutorial018-Pandas-Advanced-WorkingWithData

November 18, 2022

## 1 Pandas, indexing and other advanced data manipulation features

The past few tutorials were focussed on **Pandas**. We met some of the basic data structures in **pandas**.

Basic **pandas** objects:

- Index
- Series
- Data Frame

We also learned how these three things are related. Namely, we can think of a **pandas DataFrame** as being composed of several *named columns*, each of which is like a **Series**, and a special **Index** column along the left-hand side.

This tutorial focuses on more advanced **pandas** options to accessing, addressing (indexing) and manipulating data.

### 1.0.1 Learning goals:

- advanced **pandas** objects methods – the “verbs” that make them do useful things
- indexing and accessing row/column subsets fo data
- grouped data: aggregation and pivot tables

### 1.1 Make a data frame to play with

To get started this time instead of loading data from file, we will build a little data frame and take look at it to remind ourselves of this structure. We’ll build a data frame similar to a data set mentioned in a previosu tutorial.

First, import **pandas** because of course, and **numpy** in order to simulate some data.

```
[67]: import pandas as pd
import numpy as np      # to make the simulated data
```

Now we can make the data frame. It will have 4 variables of cardiovascular data for a number of patients (the number of patiencts can be specified):

- systolic blood pressure
- diastolic blood pressure
- blood oxygenation

- pulse rate

Given that Pandas `DataFrames` have a special `index` column, we'll just use the `index` as “patient ID” instead of making a fifth variable dedicated to it.

```
[68]: num_patients = 10      # specify the number of patients
```

We will use Numpy to simulate data by choosing a mean for each variable and a standard deviation. More specifically, the systolic blood pressure will have a mean of 125 and a standard deviation of 5. The diastolic pressure will have a lower mean (80) but the same standard deviation, the blood oxygenation will have a mean of 98.5 and a smaller standard deviation of 0.3. Finally, the pulse rate will have a mean of 65 and a standard deviation of 2.

```
[69]: sys_bp = np.int64(125 + 5*np.random.randn(num_patients,))
      dia_bp = np.int64(80 + 5*np.random.randn(num_patients,))
      b_oxy = np.round(98.5 + 0.3*np.random.randn(num_patients,), 2)
      pulse = np.int64(65 + 2*np.random.randn(num_patients,))
```

We will build the data frame using a dictionary:

```
[70]: # Make a dictionary with a "key" for each variable name, and
      # the "values" being the num_patients long data vectors
      df_dict = {'systolic BP' : sys_bp,
                  'diastolic BP' : dia_bp,
                  'blood oxygenation' : b_oxy,
                  'pulse rate' : pulse
                }

      our_df = pd.DataFrame(df_dict)      # Now make a data frame out of the dictionary
```

And now let's look at it.

```
[71]: our_df
```

```
[71]:
```

	systolic BP	diastolic BP	blood oxygenation	pulse rate
0	127	77	98.97	65
1	125	88	98.59	63
2	122	80	98.23	65
3	129	79	98.54	63
4	132	77	98.54	63
5	131	81	97.95	63
6	121	86	98.84	63
7	128	82	99.05	65
8	129	86	98.50	67
9	126	71	97.70	65

Complete the following exercise.

- Use the cell below to create a dataframe with the following data:
  - 16 patients

- systolic blood pressure 10% higher than the current
- diastolic blood pressure 5% lower
- blood oxygenation 2% higher
- a 4% higher pulse rate

```
[72]: num_pts = 16

sys_bp_ex = sys_bp*1.1
dia_bp_ex = dia_bp*0.95
b_oxy_ex = b_oxy*1.02
pulse_ex = pulse*1.04

df_dict = {'systolic BP' : sys_bp_ex,
           'diastolic BP' : dia_bp_ex,
           'blood oxygenation' : b_oxy_ex,
           'pulse rate' : pulse_ex
          }

df2 = pd.DataFrame(df_dict)
df2
```

```
[72]:
```

	systolic BP	diastolic BP	blood oxygenation	pulse rate
0	139.7	73.15	100.9494	67.60
1	137.5	83.60	100.5618	65.52
2	134.2	76.00	100.1946	67.60
3	141.9	75.05	100.5108	65.52
4	145.2	73.15	100.5108	65.52
5	144.1	76.95	99.9090	65.52
6	133.1	81.70	100.8168	65.52
7	140.8	77.90	101.0310	67.60
8	141.9	81.70	100.4700	69.68
9	138.6	67.45	99.6540	67.60

Now we can see the nice structure of the `DataFrame` object. We have four columns corresponding to our measurement variables, and each row is an “observation” which, in the case, corresponds to an individual patient.

To appreciate some of the features of a pandas `DataFrame`, let’s compare it with a numpy `Array` holding the same information. (Which we can do because we’re only dealing with numbers here - one of the main features of a pandas data frame is that it can hold non-numeric information too).

```
[73]: our_array = np.transpose(np.vstack((sys_bp, dia_bp, b_oxy, pulse)))
our_array
#np.vstack?
```

```
[73]: array([[127.  ,  77.  ,  98.97,  65.  ],
            [125.  ,  88.  ,  98.59,  63.  ],
            [122.  ,  80.  ,  98.23,  65.  ],
```

```
[129. , 79. , 98.54, 63. ],
[132. , 77. , 98.54, 63. ],
[131. , 81. , 97.95, 63. ],
[121. , 86. , 98.84, 63. ],
[128. , 82. , 99.05, 65. ],
[129. , 86. , 98.5 , 67. ],
[126. , 71. , 97.7 , 65. ]])
```

Complete the following exercise.

- Explore what `.vstack` does, use the `markdown` cell below to explain what it does in your own words

It stacks the array in sequence vertically (row wise).

- Use the following code cell to show a few examples where you create a numpy array and use `vstack` to change it, explain why you use those operations as examples

```
[74]: # example 1

english = np.array(['red', 'green', 'pink'])
spanish = np.array(['rojo', 'verde', 'rosa'])

translate_df = pd.DataFrame(np.vstack((english, spanish)).T)
print(translate_df)

# example 2

symbol = np.array(['H', 'He', 'Li', 'Be', 'B', 'C'])
atomic_num = np.array([1, 2, 3, 4, 5, 6])
name = (['Hydrogen', 'Helium', 'Lithium', 'Beryllium', 'Boron', 'Carbon'])

periodic_df = pd.DataFrame(np.vstack((atomic_num, symbol, name)))
print(periodic_df)
```

```
   0    1
0  red  rojo
1 green verde
2  pink  rosa

   0    1    2    3    4    5
0   1    2    3    4    5    6
1    H   He   Li   Be   B    C
2 Hydrogen Helium Lithium Beryllium Boron Carbon
```

I chose these two examples to use `.vstack` to combine arrays to make dataframes easily.

We can see here that our array, `our_array`, contains exactly the same information as our dataframe, `our_df`. There are 3 main differences between the two:

- they have different verbs – things they know how to do
- we have more ways to access the information in a data frame

- the data frame could contain non-numeric information (e.g. gender) if we wanted

(Also notice that the data frame is just prettier when printed than the numpy array)

## 1.2 Verbs

Let's look at some verbs. Intuitively, it seems like both variables should *know* how to take a mean. Let's see.

```
[75]: our_array.mean()
```

```
[75]: 92.59774999999999
```

So the numpy array does indeed know how to take the mean of itself, but it takes the mean of the entire array by default, which is not very useful in this case. If we want the mean of each variable, we have to specify that we want the means of the columns (i.e. row-wise means).

```
[76]: our_array.mean(axis=0)
```

```
[76]: array([127.    ,  80.7   ,  98.491,  64.2   ])
```

But look what happens if we ask for the mean of our data frame:

```
[77]: our_df.mean()
```

```
[77]: systolic BP      127.000
      diastolic BP    80.700
      blood oxygenation 98.491
      pulse rate      64.200
      dtype: float64
```

Visually, that is much more organized! We have the mean of each of our variables, nicely labeled by the variable name.

Data frames can also `describe()` themselves.

```
[78]: our_df.describe()
```

```
[78]:
```

	systolic BP	diastolic BP	blood oxygenation	pulse rate
count	10.00000	10.000000	10.000000	10.000000
mean	127.00000	80.700000	98.491000	64.200000
std	3.59011	5.121849	0.430102	1.398412
min	121.00000	71.000000	97.700000	63.000000
25%	125.25000	77.500000	98.297500	63.000000
50%	127.50000	80.500000	98.540000	64.000000
75%	129.00000	85.000000	98.777500	65.000000
max	132.00000	88.000000	99.050000	67.000000

Gives us a nice summary table of the data in our data frame.

Numpy arrays don't know how to do this.

```
[79]: our_array.describe()
```

```
-----  
AttributeError                                Traceback (most recent call last)  
Cell In [79], line 1  
----> 1 our_array.describe()  
  
AttributeError: 'numpy.ndarray' object has no attribute 'describe'
```

Data frames can also make histograms and boxplots of themselves. They aren't publication quality, but super useful for getting a feel for our data.

```
[ ]: our_df.hist();
```

```
[ ]: our_df.boxplot();
```

For a complete listing of what our data frame knows how to do, we can type `our_df.` and then hit the tab key.

```
[ ]: our_df.
```

Complete the following exercise.

- Use the next cell to report and describe two methods of `our_df`, explain why you chose those two.

One method was `.append` which appends rows to the end of the caller, returning a new object. Another method is `.asfreq` which converts time series to specified frequency. It returns the original data conformed to a new index with the specified frequency.

Let's return to the `mean()` function, and see what, exactly, it is returning. We can do this by assigning the output to a variable and looking at its type.

```
[ ]: our_means = our_df.mean()  
our_means
```

```
[ ]: type(our_means)
```

So it is a pandas series, but, rather than the index being 0, 1, 2, 3, the *index values are actually the names of our variables*.

If we want the mean pulse rate, *we can actually ask for it by name!*

```
[ ]: our_means['pulse rate']
```

This introduces another key feature of pandas: **you can access data by name**.

Complete the following exercise.

- Use the cell below to return the diastolic blood pressure from `our_means`

```
[ ]: our_means['diastolic BP']
```

### 1.3 Accessing data

Accessing data by name is kind of a big deal. It makes code more readable and faster and easier to write.

So, for example, let's say we wanted the mean pulse rate for our patients. Using numpy, we would have to remember or figure out which column of our numpy array was pulse rate. And we'd have to remember that Python indexes start at 0. *And* we'd have to remember that we have to tell numpy to take the mean down the columns explicitly. Ha.

So our code might look something like...

```
[ ]: np_style_means = our_array.mean(axis = 0)
pulse_mean = np_style_means[3]
pulse_mean
```

Compare that to doing it the pandas way:

```
[ ]: our_means = our_df.mean()
our_means['pulse rate']
```

The pandas way makes it very clear what we are doing! People like things to have names and, in pandas, things have names.

[Complete the following exercise.](#)

- Use the cell below to compute the mean of the diastolic pressure both using the numpy method and the pandas method:

```
[ ]: # numpy
dia_bp_mean = np_style_means[1]
print(dia_bp_mean)

# pandas
our_means['diastolic BP']
```

#### 1.3.1 Accessing data using square brackets

Let's look at our little data frame again.

```
[ ]: our_df
```

We can grab a column (variable) by name if we want:

```
[80]: our_df['pulse rate']
```

```
[80]: 0    65
      1    63
```

```

2    65
3    63
4    63
5    63
6    63
7    65
8    67
9    65
Name: pulse rate, dtype: int64

```

Doing this creates another DataFrame (or Series), so it knows how to do stuff to. This allows us to do things like, for example, compute the mean pulse rate in one step instead of two. Like this:

```
[81]: our_df['pulse rate'].mean()    # creates a series, then makes it compute its own
      ↪mean
```

```
[81]: 64.2
```

We can grab as many columns as we want by using a list of column names.

```
[82]: needed_cols = ['diastolic BP', 'systolic BP']    # make a list
      our_df[needed_cols]                               # use the list to grab columns
```

```
[82]:
```

	diastolic BP	systolic BP
0	77	127
1	88	125
2	80	122
3	79	129
4	77	132
5	81	131
6	86	121
7	82	128
8	86	129
9	71	126

We could also do this in one step.

```
[83]: our_df[['diastolic BP', 'systolic BP']]    # the inner brackets define our list
```

```
[83]:
```

	diastolic BP	systolic BP
0	77	127
1	88	125
2	80	122
3	79	129
4	77	132
5	81	131
6	86	121
7	82	128



8	86	129
9	71	126

(although the double brackets might look a little confusing at first)

Complete the following exercise.

- Use the cell below to extract blood oxygenation and pulse rate using a single line of code

```
[84]: our_df[['blood oxygenation', 'pulse rate']]
```

```
[84]:
```

	blood oxygenation	pulse rate
0	98.97	65
1	98.59	63
2	98.23	65
3	98.54	63
4	98.54	63
5	97.95	63
6	98.84	63
7	99.05	65
8	98.50	67
9	97.70	65

### 1.3.2 Getting row and row/column combinations of data: “indexing”

**Terminology Warning!** “Indexing” is a general term which means “accessing data by location”. In pandas, as we have seen, objects like DataFrames also have an “index” which is a special column of row identifiers. So, in pandas, we can index data using column names, row names (indexing using the index), or both. (We can also index into pandas data frames as if they were numpy arrays, which sometimes comes in handy.)

**Changing the index to make (row) indexing more intuitive** Speaking of indexes, it’s a little weird to have our patient IDs start at “0”. Both because “patient zero” has a special meaning and also because it’s just not intuitive to number a sequence of actual things starting at “0”.

Fortunately, pandas `DataFrame` (and `Series`) objects allow you to customize their index column fairly easily.

Let’s set the index to start at 1 rather than 0:

```
[85]: my_ind = np.linspace(1, 10, 10) # make a sequence from 1 to 10
      my_ind = np.int64(my_ind)      # change it from decimal to integer (not
      ↪really necessary, but...)
```

Let’s take a look at this index:

```
[86]: print(my_ind)
```

```
[ 1  2  3  4  5  6  7  8  9 10]
```

```
[87]: our_df.index = my_ind
```

```
[88]: our_df
```

```
[88]:
```

	systolic BP	diastolic BP	blood oxygenation	pulse rate
1	127	77	98.97	65
2	125	88	98.59	63
3	122	80	98.23	65
4	129	79	98.54	63
5	132	77	98.54	63
6	131	81	97.95	63
7	121	86	98.84	63
8	128	82	99.05	65
9	129	86	98.50	67
10	126	71	97.70	65

Complete the following exercise.

- Use the next cell to create a new index variable using numpy the variable should start at 5 and continue to 15 with 10 steps in between

```
[89]: new_ind = np.linspace(5,15,10)
new_ind = np.int64(new_ind)

print(new_ind)
```

```
[ 5  6  7  8  9 10 11 12 13 15]
```

### 1.3.3 Accessing data using `pd.DataFrame.loc[]`

In the section above, we saw that you can get columns of data out of a data frame using square brackets `[]`. Pandas data frames also know how to give you subsets of rows or row/column combinations.

The primary method for accessing specific bits of data from a pandas data frame is with the `loc[]` verb. It provides an easy way to get rows of data based upon the index column. In other words, `loc[]` is the way we use the data frame index as an index!

So this will give us the data for patient number 3:

```
[90]: our_df.loc[3]
```

```
[90]:
```

systolic BP	122.00
diastolic BP	80.00
blood oxygenation	98.23
pulse rate	65.00

Name: 3, dtype: float64

**Note!** The above call did **not** behave like a Python or numpy index! If it had, we would have gotten the data for patient number 4 because Python and numpy use *zero based indexing*.

But using the `loc[]` function gives us back the row “named” 3. We literally get what we asked for! Yay!

We can also *slice* out rows in chunks:

```
[91]: our_df.loc[3:6]
```

```
[91]:
```

	systolic BP	diastolic BP	blood oxygenation	pulse rate
3	122	80	98.23	65
4	129	79	98.54	63
5	132	77	98.54	63
6	131	81	97.95	63

Which, again, gives us what we asked for without having to worry about the zero-based business.

But `.loc[]` also allows us to get specific columns too. Like:

```
[92]: our_df.loc[3:6, 'blood oxygenation']
```

```
[92]:
```

3	98.23
4	98.54
5	98.54
6	97.95

Name: blood oxygenation, dtype: float64

For a single column, or:

```
[93]: our_df.loc[3:6, 'systolic BP': 'blood oxygenation']
```

```
[93]:
```

	systolic BP	diastolic BP	blood oxygenation
3	122	80	98.23
4	129	79	98.54
5	132	77	98.54
6	131	81	97.95

for multiple columns.

In summary, there are 3 main ways to get chunks of data out of a data frame “by name”.

- square brackets (only) gives us columns, e.g. `our_df['systolic BP']`
- `loc[]` with one argument gives us rows, e.g. `our_df.loc[3]`
- `loc[]` with two arguments gives us row-column combinations, e.g. `our_df.loc[3, 'systolic BP']`

Additionally, with `loc[]`, we can specify index ranges for the rows or columns or both, e.g. `new_df.loc[3:6, 'systolic BP': 'blood oxygenation']`

One final thing about using `loc[]` is that the index column in a `DataFrame` doesn't have to be numbers. It can be date/time strings (as we'll see later on), or just plain strings (as we've seen above with `Series` objects).

[Complete the following exercise.](#)

- Use the next cell to create a data frame of heart measurements where the index is the name of the patients (name and surname, make them up!):

```
[94]: names = np.array(['John Doe', 'Jane Smith', 'Patrick Wu', 'Jenny Romeo',
↳ 'Richard Martin',
        'John Green', 'Hank Green', 'David Valorz', 'Domenica
↳ Aburto', 'Marietta Aburto'])

df2 = our_df.copy()
df2.index = names
df2
```

```
[94]:
```

	systolic BP	diastolic BP	blood oxygenation	pulse rate
John Doe	127	77	98.97	65
Jane Smith	125	88	98.59	63
Patrick Wu	122	80	98.23	65
Jenny Romeo	129	79	98.54	63
Richard Martin	132	77	98.54	63
John Green	131	81	97.95	63
Hank Green	121	86	98.84	63
David Valorz	128	82	99.05	65
Domenica Aburto	129	86	98.50	67
Marietta Aburto	126	71	97.70	65

Let's look at a summary of our data using the `describe()` method:

```
[95]: our_sum = our_df.describe()
our_sum
```

```
[95]:
```

	systolic BP	diastolic BP	blood oxygenation	pulse rate
count	10.00000	10.000000	10.000000	10.000000
mean	127.00000	80.700000	98.491000	64.200000
std	3.59011	5.121849	0.430102	1.398412
min	121.00000	71.000000	97.700000	63.000000
25%	125.25000	77.500000	98.297500	63.000000
50%	127.50000	80.500000	98.540000	64.000000
75%	129.00000	85.000000	98.777500	65.000000
max	132.00000	88.000000	99.050000	67.000000

This looks suspiciously like a data frame except the index column looks like they're... er... not indexes. Let's see.

```
[96]: type(our_sum)
```

```
[96]: pandas.core.frame.DataFrame
```

Yep, it's a data frame! But let's see if that index column actually works:

```
[100]: our_sum.loc['mean']
```

```
[100]: systolic BP      127.000
       diastolic BP     80.700
       blood oxygenation 98.491
       pulse rate       64.200
       Name: mean, dtype: float64
```

Note that, with a `Series` object, we use square brackets (only) to get rows. With a `DataFrame`, square brackets (only) are used to get columns. It won't work for `DataFrame` objects:

```
[101]: our_sum['mean']
```

```
-----
KeyError                                Traceback (most recent call last)
File /opt/homebrew/lib/python3.10/site-packages/pandas/core/indexes/base.py:
  3800, in Index.get_loc(self, key, method, tolerance)
    3799 try:
-> 3800     return self._engine.get_loc(casted_key)
    3801 except KeyError as err:

File /opt/homebrew/lib/python3.10/site-packages/pandas/_libs/index.pyx:138, in
  pandas._libs.index.IndexEngine.get_loc()

File /opt/homebrew/lib/python3.10/site-packages/pandas/_libs/index.pyx:165, in
  pandas._libs.index.IndexEngine.get_loc()

File pandas/_libs/hashtable_class_helper.pxi:5745, in pandas._libs.hashtable.
  PyObjectHashTable.get_item()

File pandas/_libs/hashtable_class_helper.pxi:5753, in pandas._libs.hashtable.
  PyObjectHashTable.get_item()

KeyError: 'mean'
```

The above exception was the direct cause of the following exception:

```
KeyError                                Traceback (most recent call last)
Cell In [101], line 1
----> 1 our_sum['mean']

File /opt/homebrew/lib/python3.10/site-packages/pandas/core/frame.py:3805, in
  DataFrame.__getitem__(self, key)
    3803 if self.columns.nlevels > 1:
    3804     return self._getitem_multilevel(key)
-> 3805 indexer = self.columns.get_loc(key)
    3806 if is_integer(indexer):
```

```

3807     indexer = [indexer]

File /opt/homebrew/lib/python3.10/site-packages/pandas/core/indexes/base.py:
↪3802, in Index.get_loc(self, key, method, tolerance)
3800     return self._engine.get_loc(casted_key)
3801 except KeyError as err:
-> 3802     raise KeyError(key) from err
3803 except TypeError:
3804     # If we have a listlike key, _check_indexing_error will raise
3805     # InvalidIndexError. Otherwise we fall through and re-raise
3806     # the TypeError.
3807     self._check_indexing_error(key)

KeyError: 'mean'

```

So, with a `DataFrame`, we have to use `.loc[]` to get rows.

And now we can slice out (get a range of) rows:

```
[102]: our_sum.loc['count':'std']
```

```
[102]:
```

	systolic BP	diastolic BP	blood oxygenation	pulse rate
count	10.00000	10.000000	10.000000	10.000000
mean	127.00000	80.700000	98.491000	64.200000
std	3.59011	5.121849	0.430102	1.398412

Or rows and columns:

```
[103]: our_sum.loc['count':'std', 'systolic BP':'diastolic BP']
```

```
[103]:
```

	systolic BP	diastolic BP
count	10.00000	10.000000
mean	127.00000	80.700000
std	3.59011	5.121849

### 1.3.4 Accessing data using `pd.DataFrame.iloc[]`

Occasionally, you might want to treat a pandas `DataFrame` as a numpy `Array` and index into it using the *implicit* row and column indexes (which start as zero of course). So support this, pandas `DataFrame` objects also have an `iloc[]`.

Let's look at our data frame again:

```
[104]: our_df
```

```
[104]:
```

	systolic BP	diastolic BP	blood oxygenation	pulse rate
1	127	77	98.97	65
2	125	88	98.59	63
3	122	80	98.23	65

4	129	79	98.54	63
5	132	77	98.54	63
6	131	81	97.95	63
7	121	86	98.84	63
8	128	82	99.05	65
9	129	86	98.50	67
10	126	71	97.70	65

And let's check its shape:

```
[105]: our_df.shape
```

```
[105]: (10, 4)
```

At some level, then, Python considers this to be just a 10x4 array (like a numpy array). This is where `iloc[]` comes in; `iloc[]` will treat the data frame as though it were a numpy array – no names!

So let's index into our-df using `iloc[]`:

```
[106]: our_df.iloc[3]  # get the fourth row
```

```
[106]: systolic BP      129.00
       diastolic BP      79.00
       blood oxygenation  98.54
       pulse rate        63.00
       Name: 4, dtype: float64
```

And compare that to using `loc[]`:

```
[107]: our_df.loc[3]
```

```
[107]: systolic BP      122.00
       diastolic BP      80.00
       blood oxygenation  98.23
       pulse rate        65.00
       Name: 3, dtype: float64
```

And of course you can slice out rows and columns:

```
[108]: our_df.iloc[2:5, 0:2]
```

```
[108]:   systolic BP  diastolic BP
       3         122           80
       4         129           79
       5         132           77
```

Indexing using `iloc[]` is rarely needed on regular data frames (if you're using it, you should probably be working with a numpy `Array`).

It is, however, very handy for pulling data out of summary data tables (see below).

## 1.4 Non-numerical information (categories or factors)

One of the huge benefits of pandas objects is that, unlike numpy arrays, they can contain categorical variables.

### 1.4.1 Make another data frame to play with

Let's use tools we've learned to make a data frame that has both numerical and categorical variables.

First, we'll make the numerical data:

```
[109]: num_patients = 20      # specify the number of patients

# make some simulated data with realistic numbers.
sys_bp = np.int64(125 + 5*np.random.randn(num_patients,))
dia_bp = np.int64(80 + 5*np.random.randn(num_patients,))
b_oxy = np.round(98.5 + 0.3*np.random.randn(num_patients,), 2)
pulse = np.int64(65 + 2*np.random.randn(num_patients,))
```

(Now we'll make them interesting – this will be clear later)

```
[110]: sys_bp[0:10] = sys_bp[0:10] + 15
dia_bp[0:10] = dia_bp[0:10] + 15
sys_bp[0:5] = sys_bp[0:5] + 5
dia_bp[0:5] = dia_bp[0:5] + 5
sys_bp[10:15] = sys_bp[10:15] + 5
dia_bp[10:15] = dia_bp[10:15] + 5
```

Now let's make a categorical variable indicating whether the patient is diabetic or not. We'll make the first half be diabetic.

```
[111]: diabetic = pd.Series(['yes', 'no']) # make the short series
diabetic = diabetic.repeat(num_patients/2) # repeat each over two cell's
      ↪worth of data
diabetic = diabetic.reset_index(drop=True) # reset the series's index value
```

```
[112]: print(diabetic)
```

```
0    yes
1    yes
2    yes
3    yes
4    yes
5    yes
6    yes
7    yes
8    yes
9    yes
```



```
10    no
11    no
12    no
13    no
14    no
15    no
16    no
17    no
18    no
19    no
dtype: object
```

Now will make an “inner” sex variable.

```
[113]: sex = pd.Series(['male', 'female'])           # make the short series
```

```
[114]: print(sex)
```

```
0    male
1    female
dtype: object
```

```
[115]: sex = sex.repeat(num_patients/4)           # repeat each over one cell's
      ↪worth of data
```

```
[116]: print(sex)
```

```
0    male
0    male
0    male
0    male
0    male
1    female
1    female
1    female
1    female
1    female
dtype: object
```

```
[117]: sex = pd.concat([sex]*2, ignore_index=True) # stack or "concatenate" two
      ↪copies
```

```
[118]: print(sex)
```

```
0    male
1    male
2    male
3    male
4    male
```

```
5     female
6     female
7     female
8     female
9     female
10    male
11    male
12    male
13    male
14    male
15    female
16    female
17    female
18    female
19    female
dtype: object
```

Now we'll make a dictionary containing all our data.

```
[119]: # Make a dictionary with a "key" for each variable name, and
# the "values" being the num_patients long data vectors
df_dict = {'systolic BP' : sys_bp,
           'diastolic BP' : dia_bp,
           'blood oxygenation' : b_oxy,
           'pulse rate' : pulse,
           'sex': sex,
           'diabetes': diabetic
          }
```

And turn it into a data frame.

```
[120]: new_df = pd.DataFrame(df_dict)    # Now make a data frame out of the dictionary
```

Finally, let's up our game and make a more descriptive index column!

```
[121]: basename = 'patient '           # make a "base" row name
my_index = []                         # make an empty list
for i in range(1, num_patients+1) :   # use a for loop to add
    my_index.append(basename + str(i)) # id numbers so the base name
```

Assign our new row names to the index of our data frame.

```
[122]: new_df.index = my_index
```

Let's look at our creation!

```
[123]: new_df
```

```

[123]:      systolic BP  diastolic BP  blood oxygenation  pulse rate  sex \
patient 1      149      106      98.37      66      male
patient 2      146      96      98.89      66      male
patient 3      148      93      98.74      64      male
patient 4      139      102      98.63      60      male
patient 5      144      96      98.52      67      male
patient 6      143      95      98.74      64      female
patient 7      143      89      98.27      62      female
patient 8      141      98      98.40      64      female
patient 9      136      98      98.71      63      female
patient 10     133      98      98.68      64      female
patient 11     138      86      98.21      65      male
patient 12     123      90      98.56      66      male
patient 13     125      91      98.52      65      male
patient 14     136      85      98.43      69      male
patient 15     128      81      98.43      68      male
patient 16     132      82      99.07      65      female
patient 17     120      73      98.34      69      female
patient 18     129      77      98.57      62      female
patient 19     120      82      98.52      65      female
patient 20     123      80      98.79      67      female

```

```

      diabetes
patient 1      yes
patient 2      yes
patient 3      yes
patient 4      yes
patient 5      yes
patient 6      yes
patient 7      yes
patient 8      yes
patient 9      yes
patient 10     yes
patient 11     no
patient 12     no
patient 13     no
patient 14     no
patient 15     no
patient 16     no
patient 17     no
patient 18     no
patient 19     no
patient 20     no

```

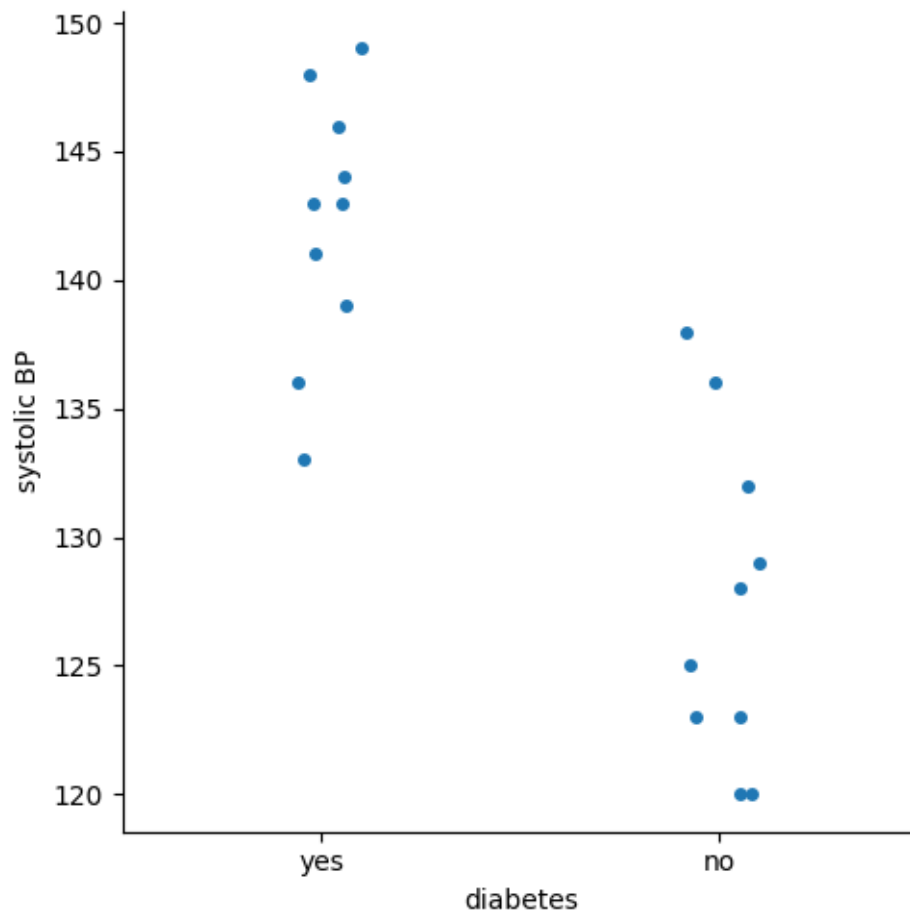
### 1.4.2 Looking at our data

Another really nice thing about pandas `DataFrames` is that they naturally lend themselves to interrogation via the visualization library `Seaborn` (we will learn about this library more in future tutorials).

So let's peek at some stuff.

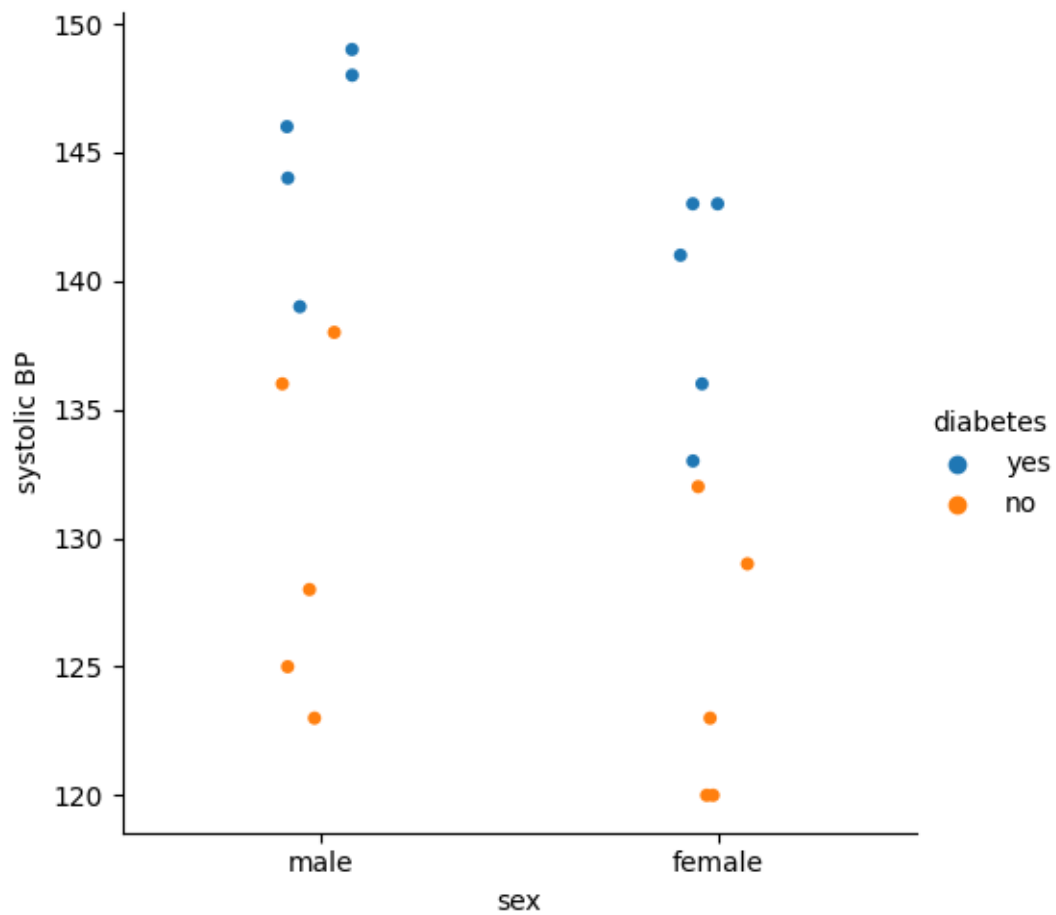
```
[124]: import seaborn as sns

sns.catplot(data=new_df, x='diabetes', y='systolic BP');
```

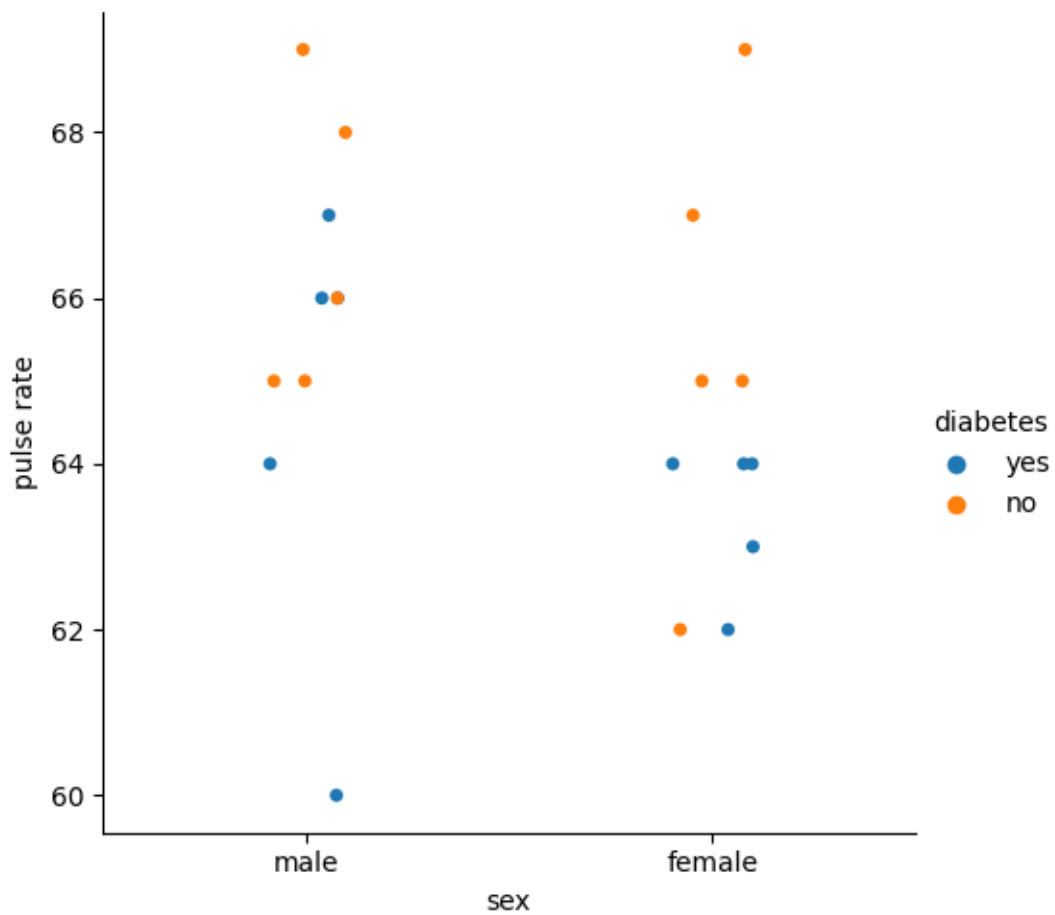


Okay, now let's go crazy and do a bunch of plots.

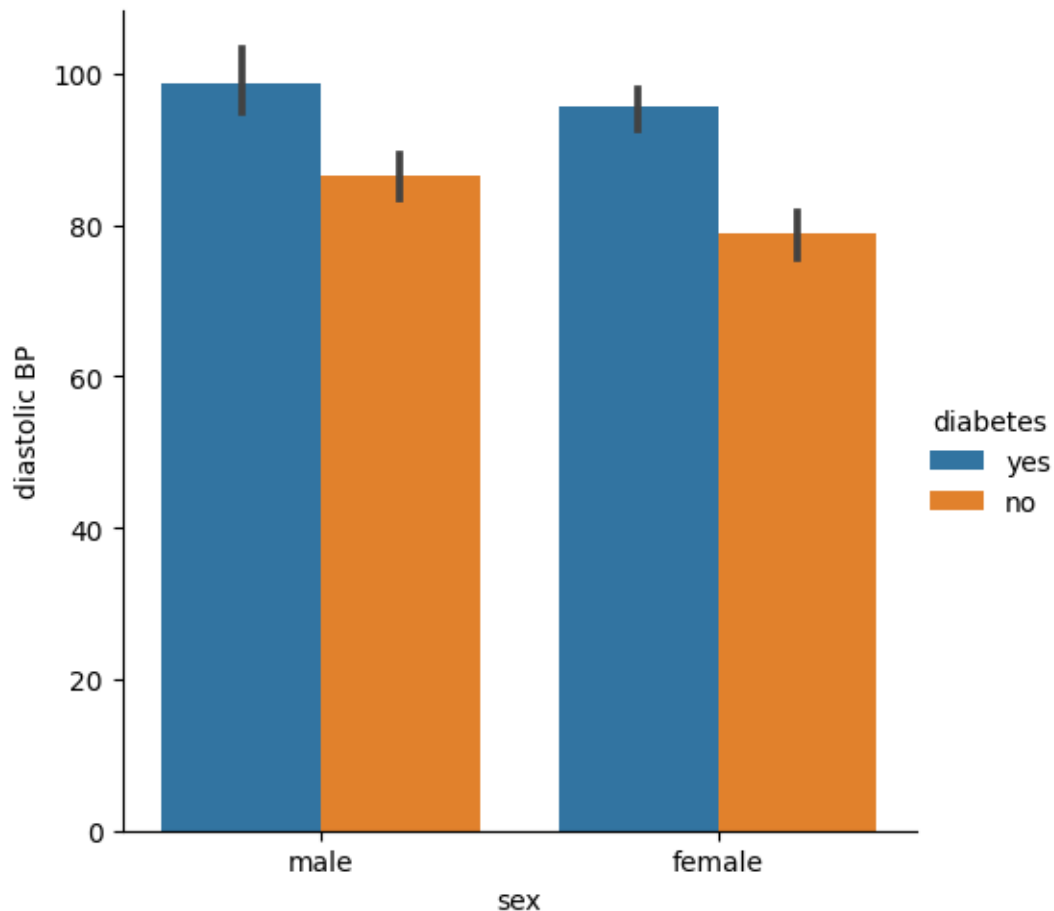
```
[125]: sns.catplot(data=new_df, x='sex', y='systolic BP', hue='diabetes');
```



```
[126]: sns.catplot(data=new_df, x='sex', y='pulse rate', hue='diabetes');
```



```
[127]: sns.catplot(data=new_df, x='sex', y='diastolic BP', hue='diabetes', kind='bar');
```



### 1.4.3 Computing within groups

Now that we have an idea of what's going on, let's look at how we could go about computing things like the mean systolic blood pressure in females vs. males, etc.

**Using the `groupby()` method** Data frames all have a `group_by()` method that, as the name implies, will group our data by a categorical variable. Let's try it.

```
[128]: new_df.groupby('sex')
```

```
[128]: <pandas.core.groupby.generic.DataFrameGroupBy object at 0x1281c2680>
```

So this gave us a `DataFrameGroupBy` object which, in and of itself, is very useful. However, *it knows how to do things!*

In general, `GroupBy` objects know how to do pretty much anything that regular `DataFrame` objects do. So, if we want the mean by gender, we can ask the `GroupBy` (for short) object to give us the mean:

```
[129]: new_df.groupby('sex').mean()
```

```
/var/folders/zc/6v283x0929j5f38j6cvlvbwr0000gn/T/ipykernel_10605/4080839992.py:1
: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is
deprecated. In a future version, numeric_only will default to False. Either
specify numeric_only or select only columns which should be valid for the
function.
```

```
new_df.groupby('sex').mean()
```

```
[129]:      systolic BP  diastolic BP  blood oxygenation  pulse rate
sex
female      132.0      87.2      98.609      64.5
male       137.6      92.6      98.530      65.6
```

**Using the `groupby()` followed by `aggregate()`** More powerfully, we can use a `GroupBy` object's `aggregate()` method to compute many things at once.

```
[130]: new_df.groupby('diabetes').aggregate(['mean', 'std', 'min', 'max'])
```

```
/var/folders/zc/6v283x0929j5f38j6cvlvbwr0000gn/T/ipykernel_10605/2935691488.py:1
: FutureWarning: ['sex'] did not aggregate successfully. If any error is raised
this will raise in a future version of pandas. Drop these columns/ops to avoid
this warning.
```

```
new_df.groupby('diabetes').aggregate(['mean', 'std', 'min', 'max'])
```

```
[130]:      systolic BP      diastolic BP      \
              mean      std  min  max      mean      std  min  max
diabetes
no      127.4  6.363088  120  138      82.7  5.538752  73   91
yes     142.2  5.094660  133  149      97.1  4.653553  89  106

      blood oxygenation      pulse rate      \
              mean      std  min  max      mean      std  min
diabetes
no      98.544  0.240009  98.21  99.07      66.1  2.183270  62
yes     98.595  0.197386  98.27  98.89      64.0  2.054805  60

      max
diabetes
no      69
yes     67
```

Okay, what's going on here? First, we got a lot of information out. Second, we got a warning because pandas couldn't compute the mean, etc., on the gender variable, which is perfectly reasonable of course.

We can handle this by using our skills to carve out a subset of our data frame – just the columns



of interest – and then use `groupby()` and `aggregate()` on that.

```
[131]: temp_df = new_df[['systolic BP', 'diastolic BP', 'diabetes']]      # make a
      ↪data frame with only the columns we want
our_summary = temp_df.groupby('diabetes').aggregate(['mean',
      ↪'std', 'min', 'max'])      ↪
      ↪# compute stuff on those columns
our_summary
```

```
[131]:
```

	systolic BP				diastolic BP			
	mean	std	min	max	mean	std	min	max
diabetes								
no	127.4	6.363088	120	138	82.7	5.538752	73	91
yes	142.2	5.094660	133	149	97.1	4.653553	89	106

Notice here that there are *groups of columns*. Like there are two “meta-columns”, each with four data columns in them. This makes getting the actual values out of the table for further computation, etc., kind of a pain. It’s called “multi-indexing” or “hierarchical indexing”. It’s a pain.

Here are a couple examples.

```
[132]: our_summary[("systolic BP", "mean")]
```

```
[132]: diabetes
no      127.4
yes      142.2
Name: (systolic BP, mean), dtype: float64
```

```
[133]: our_summary.loc[("no")]
```

```
[133]: systolic BP    mean    127.400000
      std         6.363088
      min    120.000000
      max    138.000000
diastolic BP    mean     82.700000
      std         5.538752
      min    73.000000
      max     91.000000
Name: no, dtype: float64
```

Of course, we could do the blood pressure variables separately and store them for later plotting, etc.

```
[134]: temp_df = new_df[['systolic BP', 'diabetes']]      # make a data frame with
      ↪only the columns we want
our_summary = temp_df.groupby('diabetes').aggregate(['mean',
      ↪'std', 'min', 'max'])      ↪
      ↪# compute stuff on those columns
our_summary
```

```
[134]:
```

	systolic BP			
	mean	std	min	max
diabetes				
no	127.4	6.363088	120	138
yes	142.2	5.094660	133	149

But we still have a meta-column label!

Here's where `.iloc[]` comes to the rescue!

If we look at the shape of the summary:

```
[135]: our_summary.shape
```

```
[135]: (2, 4)
```

We see that, ultimately, the data is just a 2x4 table. So if we want, say, the standard deviation of non-diabetics, we can just do:

```
[136]: our_summary.iloc[0, 1]
```

```
[136]: 6.363087999461338
```

And we get back a pure number.

We can also do things “backwards”, that is, instead of subsetting the data and then doing a `groupby()`, we can do the `groupby()` and then index into it and compute what we want. For example, if we wanted the mean of systolic blood pressure grouped by whether patients had diabetes or not, we could go one of two ways.

We could subset and then group:

```
[137]: new_df[['systolic BP', 'diabetes']].groupby('diabetes').mean()
```

```
[137]:
```

	systolic BP
diabetes	
no	127.4
yes	142.2

Or we could group and then subset:

```
[138]: new_df.groupby('diabetes')[['systolic BP']].mean()
```

```
[138]:
```

	systolic BP
diabetes	
no	127.4
yes	142.2

Okay, first, it's cool that there are multiple ways to do things. Second – **aarrgghh!** – things are starting to get complicated and code is getting hard to read!

**Using pivot tables** “Pivot tables” (so named because allow you to look at data along different dimensions or directions) provide a handy solution for summarizing data.

By default, pivot tables tabulate the mean of data. So if we wish to compute the average systolic blood pressure broken out by diabetes status, all we have to do is:

```
[139]: new_df.pivot_table('systolic BP', index='diabetes')
```

```
[139]:          systolic BP
diabetes
no           127.4
yes          142.2
```

Here, `index` is used in the “row names” sense of the word.

We can also have another grouping variables map to the columns of the output if we wish:

```
[140]: new_df.pivot_table('systolic BP', index='diabetes', columns='sex')
```

```
[140]: sex      female    male
diabetes
no       124.8    130.0
yes      139.2    145.2
```

Finally, we can specify pretty much any other summary function we want to “aggregate” by:

```
[141]: new_df.pivot_table('systolic BP', index='diabetes', columns='sex',
    ↪aggfunc='median')
```

```
[141]: sex      female    male
diabetes
no       123      128
yes      141      146
```

If you want to customize the column names using the aggregate function, you can (Though it is somewhat limited)! Look at the example down below for an explanation

```
[142]: new_df.groupby('diabetes').aggregate(Mean=('systolic BP', "mean"))
```

```
[142]:          Mean
diabetes
no       127.4
yes      142.2
```

The “Mean” is your new title, while inside the second set of parantheses is where/what you want the aggregate function to calculate

However, as you might have noticed, this is fairly limited. It removes the meta column titles, replacing them with the title of your choice. This can make it somewhat difficult to interpret your tables. Additionally, you can’t have any spaces in the new title of your choice.

```
[143]: new_df.groupby('diabetes').aggregate(Mean=('systolic BP',"mean"),
                                           Standard_Deviation = ('systolic BP',
                                           ↪"std"))
```

```
[143]:
```

	Mean	Standard_Deviation
diabetes		
no	127.4	6.363088
yes	142.2	5.094660

## 1.5 vs.

```
[144]: new_df.groupby('diabetes').aggregate( Mean=('systolic BP',"mean"), STD = ↪
↪('systolic BP',"std"))
```

```
[144]:
```

	Mean	STD
diabetes		
no	127.4	6.363088
yes	142.2	5.094660

(Where `aggfunc` can me ‘min’, ‘sum’, ‘std’, etc., etc.)

## 1.6 Summary

In this tutorial, we have covered some key aspects of working with data using pandas data frames. These were:

- doing things with data using the methods – the verbs – of pandas objects
- accessing subsets of the data with
  - square brackets
  - the `.loc[]` method
  - the `.iloc[]` method
- assembling data frames and customizing the index
- grouping data and computing summaries using
  - `groupby()` and `aggregate()`
  - pivot tables

## 1.7 Complete the following exercise.

1. Make a data frame that has
  - one categorical variable, “bilingual”, that splits the data in half (“yes” and “no”)
  - two numerical variables, verbal GRE and quant GRE
  - (you can build in, or not, whatever effect of bilingual you wish)
  - (GRE scores have a mean of about 151 and a std. dev. of about 8.5)
2. Set the index to be “Student 1”, “Student 2”, etc.
3. Do a seaborn plot of verbal GRE vs. bilinguality (is that a word?)
4. Make another one of quant GRE vs. bilingual status
5. Compute the mean and standard *error* of each score separated by bilingual status (using any method you wish!)

```
[151]: num_participant = 20

# bilingual array
bilingual = pd.Series(['yes', 'no'])
bilingual = bilingual.repeat(num_participant/2)
bilingual = bilingual.reset_index(drop=True)

# verbal and quant GRE
verbal = np.int64(151 + 8.5 * np.random.randn(num_participant,))
quant = np.int64(151 + 8.5 * np.random.randn(num_participant,))

verbal[0:10] = verbal[0:10] + 6
quant[0:10] = quant[0:10] + 3

# make dict
GRE_bi = {'Verbal GRE': verbal,
          'Quant. GRE': quant,
          'Bilingual': bilingual,
          }

# dataframe
GRE_bi_df = pd.DataFrame(GRE_bi)

# index students
index = np.array([])
for i in range(1, num_participant + 1):
    index = np.append(index, "Student " + str(i))

GRE_bi_df.index = index # specify index
GRE_bi_df
```

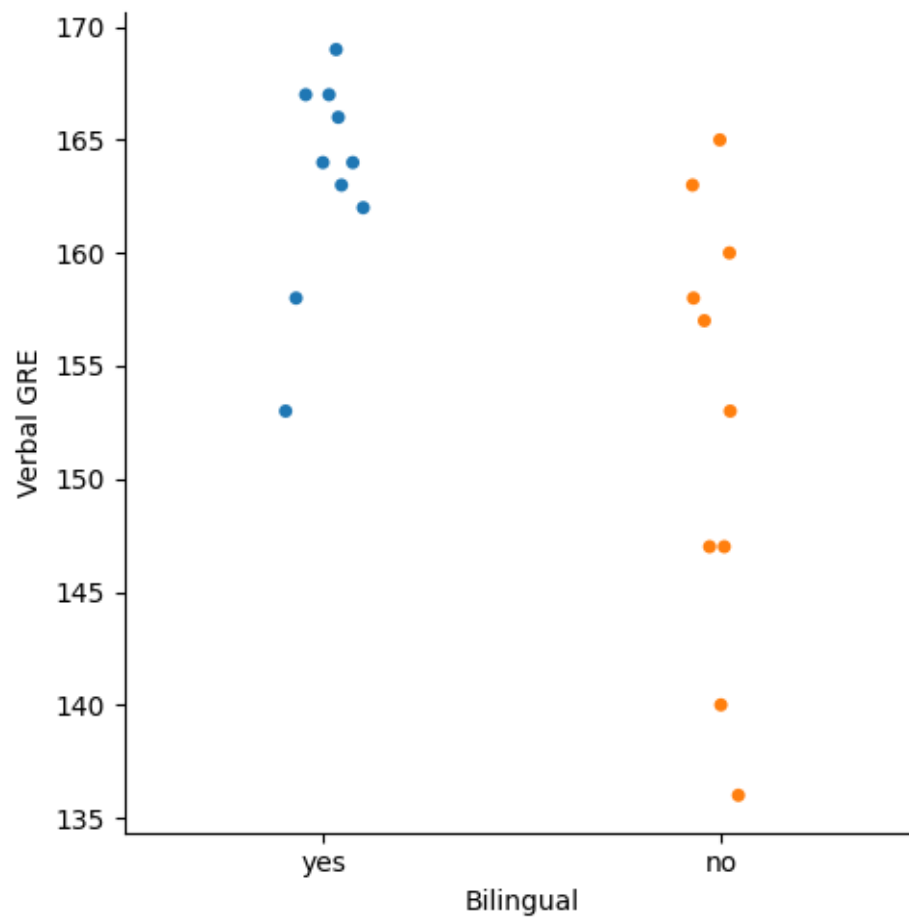
```
[151]:
```

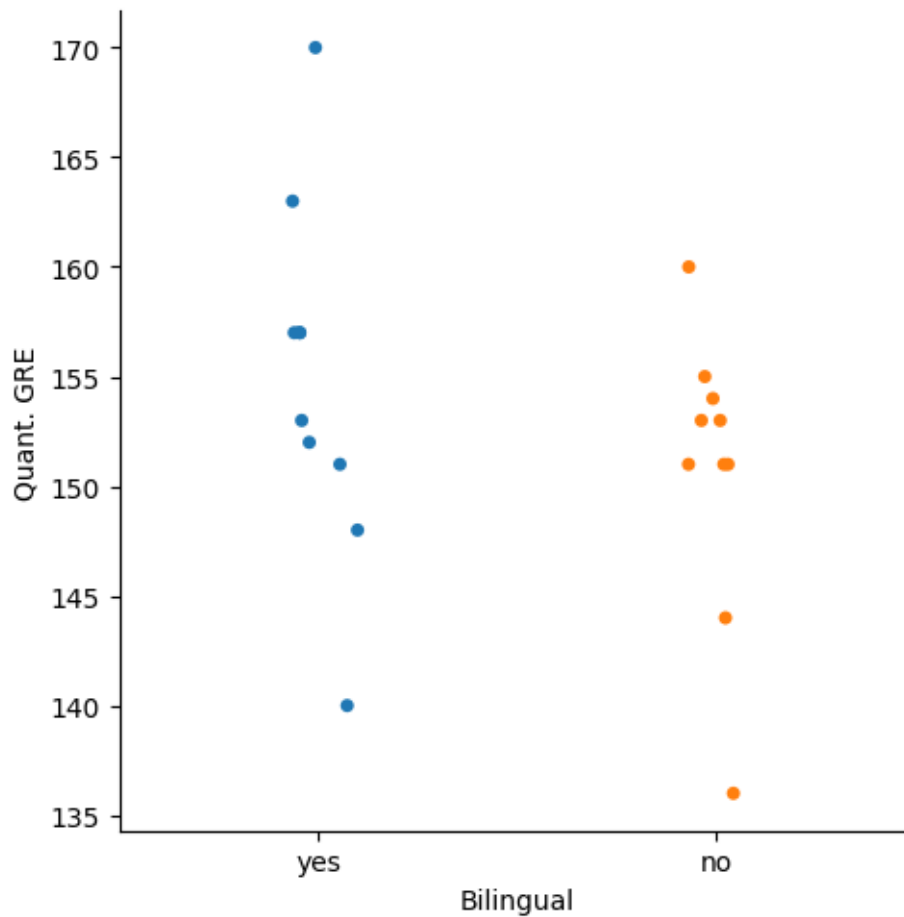
	Verbal GRE	Quant. GRE	Bilingual
Student 1	166	152	yes
Student 2	164	157	yes
Student 3	167	157	yes
Student 4	158	170	yes
Student 5	162	148	yes
Student 6	163	140	yes
Student 7	153	151	yes
Student 8	167	157	yes
Student 9	164	163	yes
Student 10	169	153	yes
Student 11	147	136	no
Student 12	153	144	no
Student 13	140	153	no
Student 14	160	153	no
Student 15	136	151	no

Student 16	158	155	no
Student 17	165	151	no
Student 18	163	151	no
Student 19	147	154	no
Student 20	157	160	no

```
[153]: # verbal GRE vs. bilingual
sns.catplot(data = GRE_bi_df, x = 'Bilingual', y = 'Verbal GRE', hue = 'Bilingual');

# quant GRE vs. bilingual
sns.catplot(data = GRE_bi_df, x = 'Bilingual', y = 'Quant. GRE', hue = 'Bilingual');
```





```
[158]: # mean and standard error of each score separated by bilingual status
GRE_sum = GRE_bi_df.groupby('Bilingual').aggregate(['mean', 'std'])
GRE_sum
```

```
[158]:
```

	Verbal GRE		Quant. GRE	
Bilingual	mean	std	mean	std
no	152.6	9.788883	150.8	6.562520
yes	163.3	4.762119	154.8	8.189424